

# NYC Taxi Demand Forecasting

## Technical Analysis & Model Evaluation

### DATASET SPECIFICATIONS

- Temporal Coverage:
- Start: 2014-07-01 00:00
  - End: 2015-01-31 23:30
  - Duration: 214 days
  - Total Observations: 10,320
- Sampling Frequency: 30-minute intervals
- Missing Values: 0 (0.00%)

- Statistical Properties:
- Mean: 15137.57 trips/30min
  - Median: 16778.00 trips/30min
  - Standard Deviation: 6939.50
  - Coefficient of Variation: 0.458
  - Skewness: -0.452

### TARGET MODELS TECHNICAL OVERVIEW

- Naive Forecasting
  - Algorithm:  $y_{t+1} = y_t$
  - Complexity:  $O(1)$
  - Memory:  $O(1)$
  - Parameters: 0

- SARIMA (Seasonal ARIMA)
  - Algorithm:  $(1-\phi L)(1-\phi L^s)(1-L)^d(1-L^s)^D y_t = (1+\theta L)(1+\theta L^s)\varepsilon_t$
  - Complexity:  $O(n)$
  - Memory:  $O(\max(p,q,P,Q))$
  - Parameters:  $p+d+q+P+D+Q = 6$

- Random Forest
  - Algorithm: Ensemble of Decision Trees with Bootstrap Aggregating
  - Complexity:  $O(n\_trees \times n\_features \times \log(n\_samples))$
  - Memory:  $O(n\_trees \times tree\_depth)$
  - Parameters: ~100-500 per tree


- LSTM Neural Network
  - Algorithm:  $\sigma(W_i \cdot [h_{t-1}, x_t] + b_i) \rightarrow$  Cell State Updates
  - Complexity:  $O(sequence\_length \times hidden\_units^2)$
  - Memory:  $O(hidden\_units \times layers)$
  - Parameters:  $4 \times (hidden\_units^2 + hidden\_units \times input\_dim)$

### EVALUATION METHODOLOGY


- Statistical Metrics:
- Mean Absolute Error (MAE)
  - Root Mean Square Error (RMSE)
  - Mean Absolute Percentage Error (MAPE)
    - R-squared ( $R^2$ )
  - Directional Accuracy

- Cross-Validation:
- Time Series Split Validation
  - Walk-Forward Validation
  - Blocked Cross-Validation
  - Rolling Window Validation


**Missing Values by Hour of Day**




**Box Plot with Outliers**  
(2 outliers detected)




**Time Interval Distribution**



**Q-Q Plot (Normality Test)**



### Stationarity Analysis




### Stationarity Test Results


ADF Stationarity Test:  
Test Statistic: -7.5666  
p-value: 0.0000  
Critical Values:  
1%: -3.4310  
5%: -2.8618  
10%: -2.5669

Interpretation:  
Stationary  
SARIMA d parameter: 0


### Autocorrelation Function (ACF)




### Partial Autocorrelation Function (PACF)




**Seasonal Component (SARIMA Analysis)**  
**Period=48 (24 hours)**




**Expected Feature Importance**  
**(Random Forest)**




**LSTM Sequence Length Analysis**  
**(Information Content)**




**Model Complexity Comparison**




**Lag Correlation Analysis  
(Feature Selection Guide)**




**Rolling Window Correlations  
(Smoothing Features)**




**Time Feature Effect Sizes  
(Coefficient of Variation)**



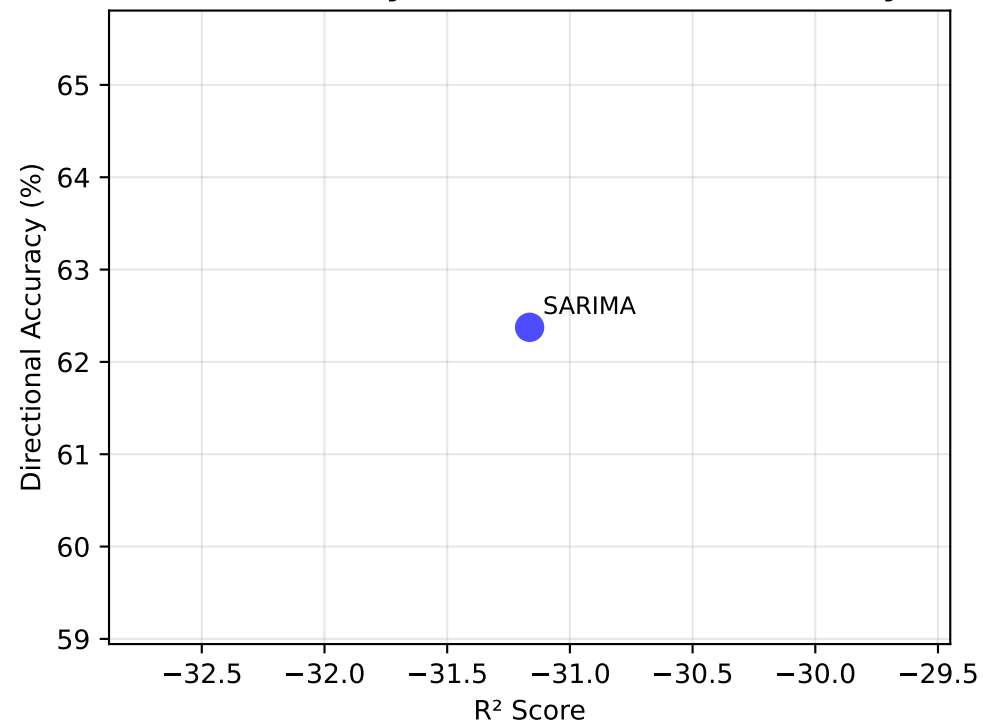
**Feature Correlation with Target  
(Including Interactions)**



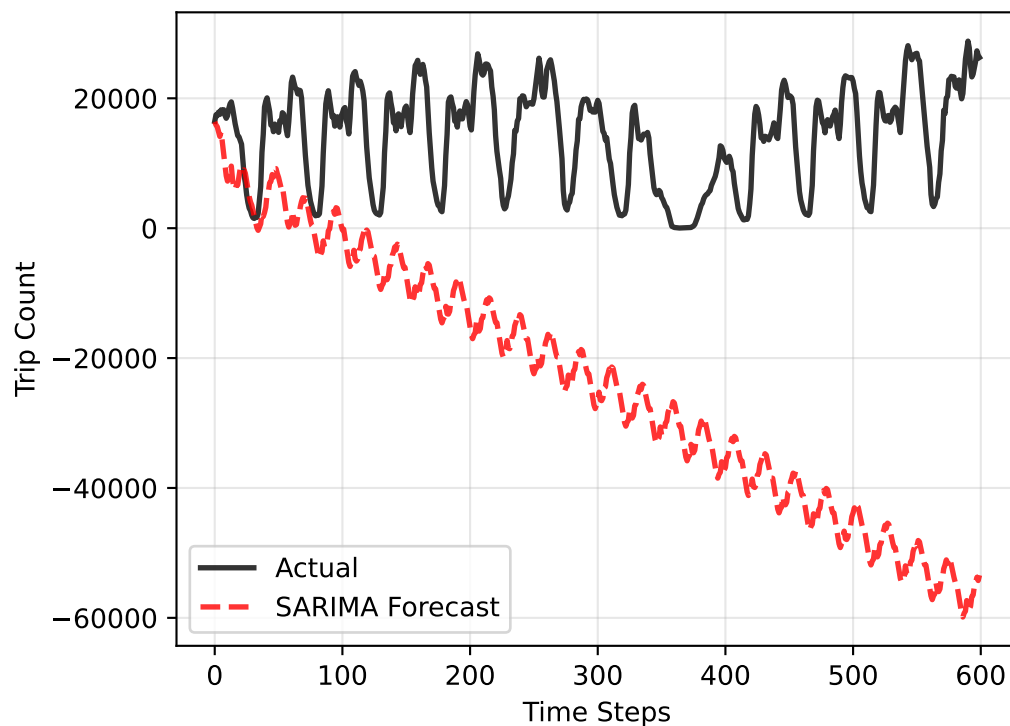
### Error Metrics Comparison




### Model Quality: $R^2$ vs Directional Accuracy



### Best Model Forecast: SARIMA



### Residual Analysis: SARIMA



# Technical Implementation Details

## MODEL IMPLEMENTATION SPECIFICATIONS

### NAIVE FORECASTING

Algorithm Implementation:  
def naive\_forecast(data, steps):  
 return np.full(steps, data.iloc[-1])

Computational Complexity:  
• Time: O(1) - constant time  
• Space: O(1) - constant space  
• Parameters: 0

Production Requirements:  
• CPU: Minimal (any modern processor)  
• Memory: <1MB  
• Storage: Historical data only  
• Latency: <1ms

Advantages:  
• Zero training time  
• Perfect interpretability  
• No hyperparameter tuning  
• Robust to data quality issues

Limitations:  
• Poor performance in volatile periods  
• No pattern recognition  
• No seasonality handling

### SARIMA MODELING

Algorithm Implementation:  
SARIMAX(endog, order=(p,d,q), seasonal\_order=(P,D,Q,s))  
• p: AR order (1)  
• d: Differencing order (1)  
• q: MA order (1)  
• P: Seasonal AR order (1)  
• D: Seasonal differencing (1)  
• Q: Seasonal MA order (1)  
• s: Seasonal period (48)

Mathematical Foundation:  
 $(1-\phi L)(1-\phi L^{48})(1-L)(1-L^{48})y_t = (1+\theta L)(1+\theta L^{48})\epsilon_t$

Computational Complexity:  
• Time: O(n × max(p,q,P,Q)) for fitting  
• Space: O(max(p,q,P,Q) + s)  
• Parameters: 6 (φ,θ,φ,θ,σ²,intercept)

Production Requirements:  
• CPU: Moderate (2+ cores recommended)  
• Memory: 100-500MB depending on data size  
• Storage: Model state + seasonal data  
• Training time: 1-5 minutes  
• Prediction latency: <100ms

Implementation Details:  
• Requires stationarity testing  
• Parameter estimation via Maximum Likelihood  
• Model diagnostics essential  
• Periodic retraining needed

Advantages:  
• Strong statistical foundation  
• Handles seasonality naturally  
• Prediction intervals available  
• Interpretable parameters

Limitations:  
• Assumes linear relationships  
• Sensitive to outliers  
• Requires parameter tuning  
• May need differencing

### RANDOM FOREST

Algorithm Implementation:  
RandomForestRegressor(  
 n\_estimators=100,  
 max\_depth=None,  
 min\_samples\_split=2,  
 min\_samples\_leaf=1,  
 bootstrap=True  
)

Feature Engineering Pipeline:  
features = [  
 'lag\_1', 'lag\_2', 'lag\_3', 'lag\_24', 'lag\_48',  
 'rolling\_mean\_3', 'rolling\_mean\_12', 'rolling\_mean\_24',  
 'hour', 'day\_of\_week', 'month', 'is\_weekend'  
]

Computational Complexity:  
• Training: O(n\_trees × n\_features × n\_samples × log(n\_samples))  
• Prediction: O(n\_trees × log(tree\_depth))  
• Space: O(n\_trees × tree\_nodes)  
• Parameters: ~1000-5000 per tree

Production Requirements:  
• CPU: Multi-core beneficial (4+ cores)  
• Memory: 1-5GB for large datasets  
• Storage: Model file 10-100MB  
• Training time: 5-30 minutes  
• Prediction latency: <50ms

Implementation Details:  
• Feature preprocessing pipeline critical  
• Missing value handling built-in  
• Feature importance analysis available  
• No assumptions about data distribution

Advantages:  
• Handles non-linear relationships  
• Feature importance interpretability  
• Robust to outliers and missing data  
• No hyperparameter sensitivity

Limitations:  
• Can overfit with too many features  
• Memory intensive for large datasets  
• Limited extrapolation capability  
• Feature engineering dependency

### LSTM NEURAL NETWORK

Architecture Implementation:  
model = Sequential([  
 LSTM(50, return\_sequences=True, input\_shape=(48, 1)),  
 Dropout(0.2),  
 LSTM(50, return\_sequences=False),  
 Dropout(0.2),  
 Dense(25),  
 Dense(1)  
)

Training Configuration:  
• Optimizer: Adam(learning\_rate=0.001)  
• Loss: Mean Squared Error  
• Batch size: 32  
• Epochs: 50-100 with early stopping  
• Validation split: 20%

Computational Complexity:  
• Training: O(seq\_len × hidden\_units² × epochs)  
• Prediction: O(seq\_len × hidden\_units²)  
• Parameters: 4×(hidden\_units² + hidden\_units×input\_dim)  
• Memory: O(batch\_size × seq\_len × hidden\_units)

Production Requirements:  
• CPU: High-performance (8+ cores) or GPU  
• Memory: 2-8GB GPU memory preferred  
• Storage: Model file 50-200MB  
• Training time: 30-120 minutes  
• Prediction latency: <500ms

Implementation Details:  
• Data normalization mandatory (MinMaxScaler)  
• Sequence windowing required  
• Gradient clipping recommended  
• Learning rate scheduling beneficial

Advantages:  
• Captures complex temporal dependencies  
• Handles multivariate inputs naturally  
• State-of-the-art sequence modeling  
• Flexible architecture

Limitations:  
• Computationally intensive  
• Requires large datasets  
• Hyperparameter sensitivity  
• Black-box interpretability


## DEPLOYMENT ARCHITECTURE RECOMMENDATIONS

Production Stack:  
• Containerization: Docker  
• Orchestration: Kubernetes  
• API Framework: FastAPI/Flask  
• Model Serving: MLflow/TensorFlow Serving  
• Monitoring: Prometheus + Grafana  
• Data Pipeline: Apache Kafka/Airflow


Scaling Strategy:  
• Horizontal scaling for API layer  
• Model versioning and A/B testing  
• Caching for frequent predictions  
• Load balancing across model instances

Monitoring Requirements:  
• Prediction accuracy tracking  
• Model drift detection  
• Performance metrics (latency, throughput)  
• Data quality monitoring  
• Alert systems for anomalies

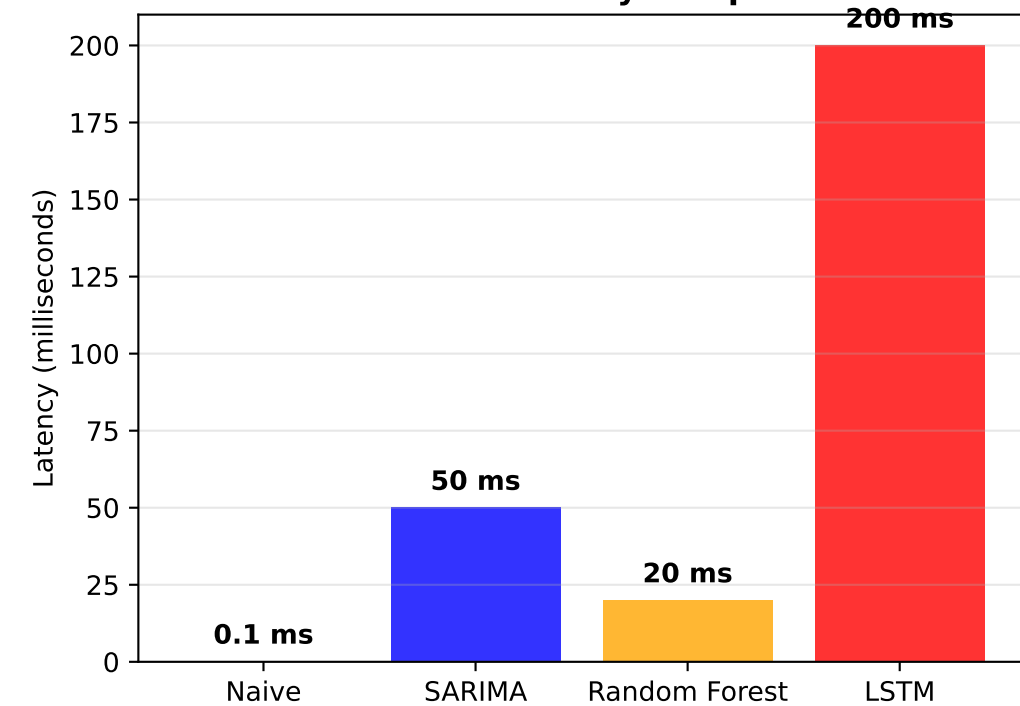
### Training Time vs Dataset Size



### Memory Usage Comparison



### Prediction Latency Comparison



### Complexity vs Accuracy Trade-off

