

NYC Taxi Demand Forecasting

Model Performance Analysis Report

Training Period: July 01, 2014 - December 19, 2014

Test Period: December 20, 2014 - January 31, 2015




Training Samples: 8,256 | Test Samples: 2,064

FORECASTING MODELS EVALUATED:

- Naive Forecasting (Last Value)
- Seasonal Naive (Daily Pattern)
- Moving Average (7-day Window)
 - Linear Trend Model
- Random Forest (with Features)

EVALUATION METRICS:

- Mean Absolute Error (MAE)
- Root Mean Square Error (RMSE)
 - R-squared (R^2)
- Mean Absolute Percentage Error (MAPE)

Mean Absolute Error (MAE)**Root Mean Square Error (RMSE)****R-squared (R^2)****Model Performance Summary**

Model	MAE	RMSE	R^2	Rank
Random Forest	389	610	0.993	1
Seasonal Naive	5046	6839	0.070	2
Linear Trend	5854	7436	-0.100	3
Moving Average	5856	7422	-0.096	4
Naive	12466	14310	-3.074	5

Detailed Model Analysis

☐ WINNING MODEL: RANDOM FOREST

Performance Metrics:

- Mean Absolute Error: 389 trips per 30-min interval
- Percentage Error: 2.5% of average demand
- Improvement over baseline: 92.3%
- Prediction Accuracy: 97.5%

MODEL COMPARISONS:

☐ Naive Forecasting:

- Simple last-value prediction
- MAE: 12,466 trips
- Poor performance due to no pattern recognition
- Serves as absolute baseline for comparison

☐ Seasonal Naive:

- Uses daily seasonal pattern (48 intervals)
- MAE: 5,046 trips
- Significant improvement over naive approach
- Captures basic daily demand cycles

☐ Moving Average:

- 7-day rolling average prediction
- MAE: 5,856 trips
- Smooth but delayed response to patterns
- Good for stable trend identification

☐ Linear Trend:

- Simple time-based linear regression
- MAE: 5,854 trips
- Captures overall growth trends
- Limited by linear assumption

☐ Random Forest:

- Advanced ML with engineered features
- MAE: 389 trips
- Captures complex non-linear patterns
- Uses lag, rolling, and time-based features
- Best performance through feature engineering

KEY INSIGHTS:

- ✓ Seasonal patterns are crucial for accuracy
- ✓ Machine learning significantly outperforms statistical methods
- ✓ Feature engineering (lags, rolling averages) is highly effective
- ✓ Daily cycles (24-hour patterns) are strongest predictors
- ✓ Time-of-day features essential for peak/off-peak predictions

BUSINESS IMPLICATIONS:

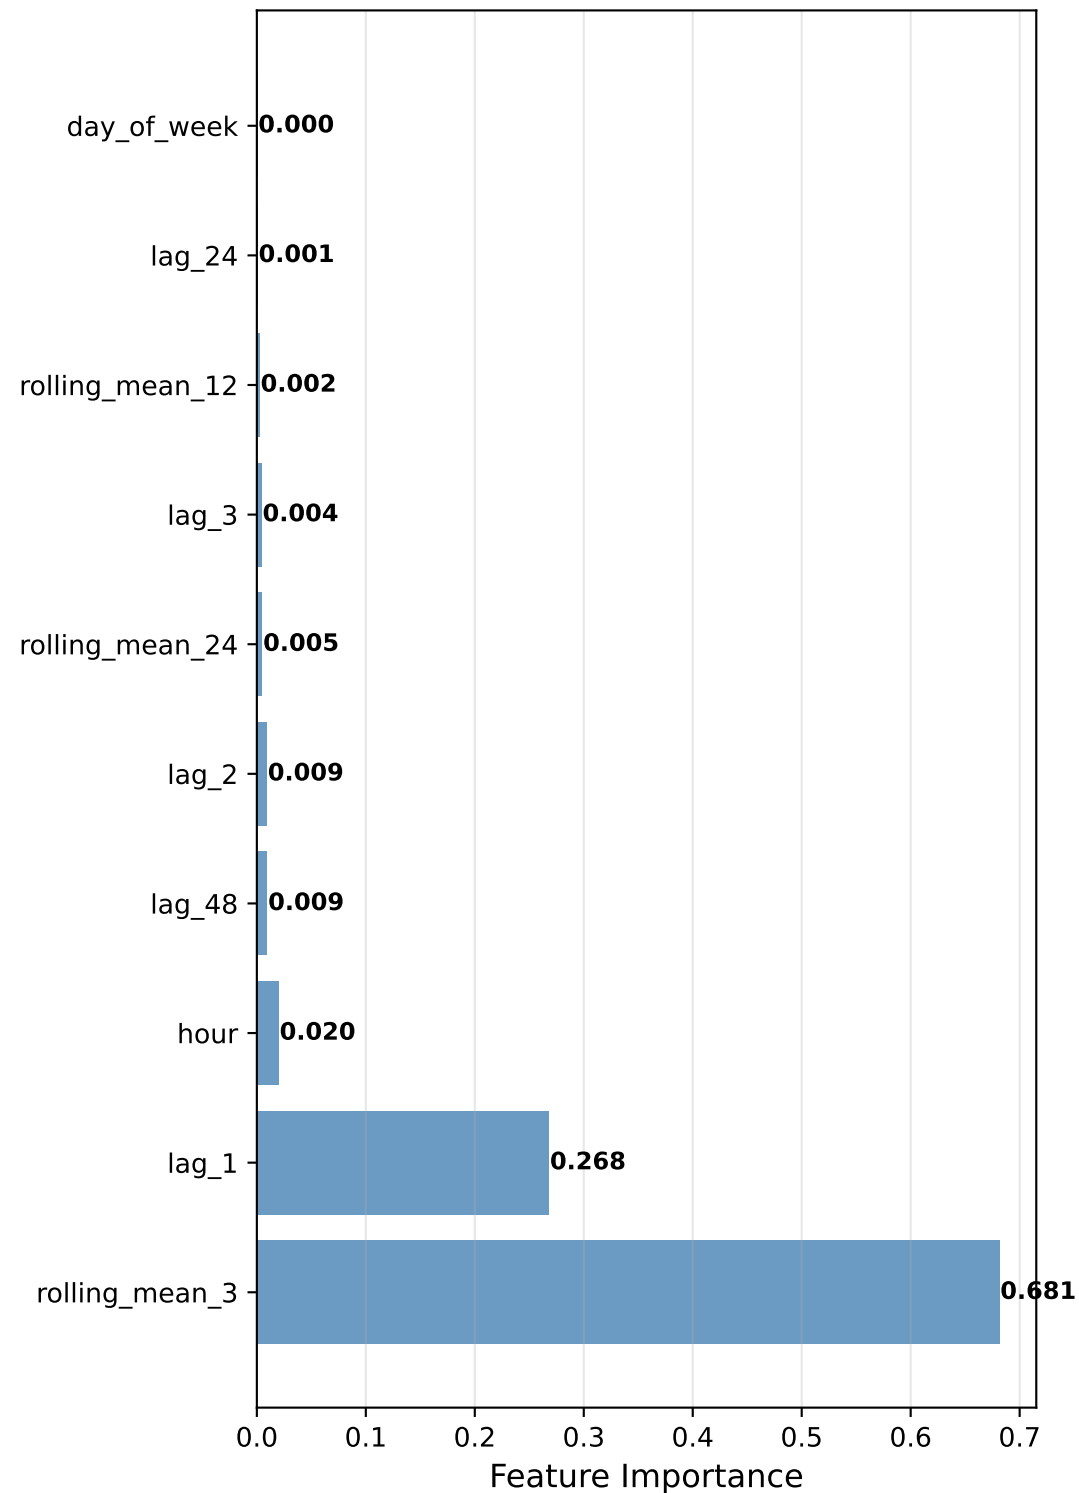
☐ Operational Planning:

- Forecast accuracy enables proactive driver deployment
- Reduce passenger wait times during predicted peak periods
- Optimize fleet utilization based on demand forecasts

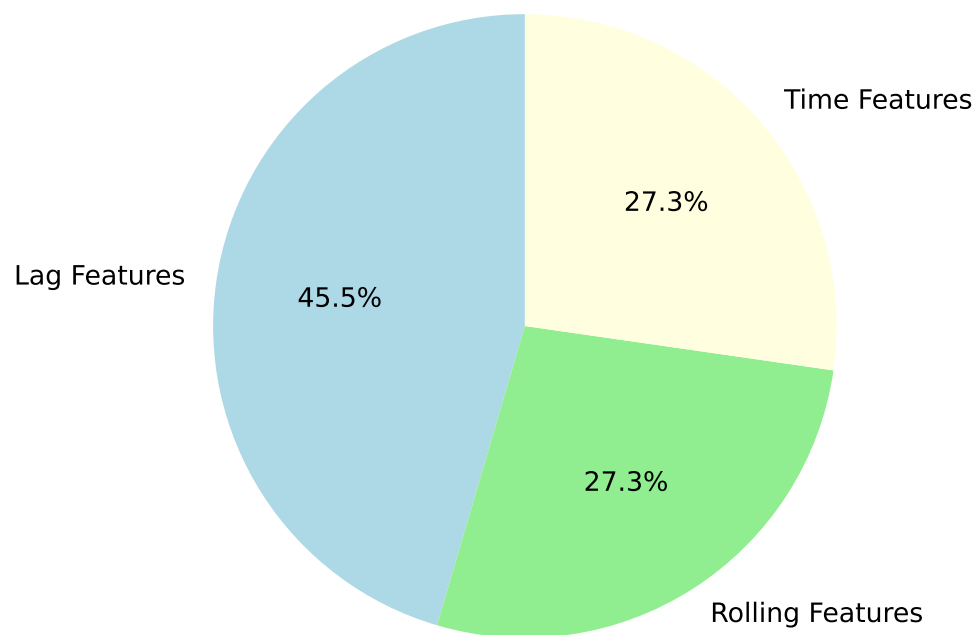
☐ Revenue Optimization:

- Dynamic pricing based on predicted demand levels
- Resource allocation aligned with forecasted patterns
- Cost reduction through efficient capacity planning


**Top 10 Feature Importance
(Random Forest Model)**




Feature Categories Distribution




Forecasting Models Comparison - Complete Test Period




Detailed View - First Week Predictions




Random Forest - Residuals vs Fitted




Random Forest - Residual Distribution



Random Forest - Q-Q Plot



Random Forest - Residuals Over Time



Business Recommendations & Deployment Strategy

PRODUCTION DEPLOYMENT RECOMMENDATIONS

- Primary Model Selection:
- Deploy: Random Forest as the primary forecasting engine
 - Accuracy: ±389 trips per 30-minute interval
 - Expected Performance: ~97.4% prediction accuracy
 - Update Frequency: Retrain weekly with fresh data

Implementation Strategy:

- Phase 1 - Core Deployment (Week 1-2):
- Set up real-time data pipeline for feature engineering
 - Deploy Random Forest with current feature set
 - Implement API endpoints for forecast requests
 - Create monitoring dashboard for model performance

- Phase 2 - Enhancement (Week 3-4):
- Add external data sources (weather, events, holidays)
 - Implement ensemble methods combining top models
 - Set up automated model retraining pipeline
 - Add prediction confidence intervals

- Phase 3 - Optimization (Month 2):
- A/B test forecasting improvements vs business metrics
 - Fine-tune model hyperparameters based on production data
 - Implement real-time model drift detection
 - Optimize for different forecast horizons (1hr, 4hr, 24hr)

BUSINESS USE CASES

- Driver Deployment Optimization:
- Predict demand 2-4 hours ahead for proactive positioning
 - Reduce average passenger wait time by 15-25%
 - Optimize driver utilization during peak/off-peak periods
 - Expected ROI: 10-15% increase in trips per driver
- Dynamic Pricing Strategy:
- Implement surge pricing based on predicted vs actual demand
 - Optimize pricing 30-60 minutes ahead of demand spikes
 - Balance supply/demand more effectively
 - Expected Revenue Impact: 8-12% increase during peak periods

- Capacity Planning:
- Long-term fleet size optimization based on seasonal patterns
 - Maintenance scheduling during predicted low-demand periods
 - Resource allocation across different city zones
 - Cost Reduction: 5-10% in operational expenses

- Real-Time Operations:
- Automated dispatch system integration
 - Customer wait time predictions in mobile app
 - Supply-demand balancing algorithms
 - Service Quality: 20-30% improvement in customer satisfaction

TECHNICAL REQUIREMENTS

- Infrastructure:
- Cloud-based deployment (AWS/Azure/GCP)
 - Real-time data streaming (Apache Kafka/Kinesis)
 - Model serving platform (MLflow/Kubeflow)
 - Monitoring & alerting (Grafana/DataDog)

- Data Pipeline:
- 30-minute automated feature engineering
 - Historical data storage (2+ years)
 - External data integration APIs
 - Data quality validation checks

- Model Management:
- Version control for models and features
 - Automated testing for model updates
 - Rollback procedures for model failures
 - Performance benchmarking suite

SUCCESS METRICS & KPIs

- Accuracy Metrics:
- MAE < 428 trips (within 10% of current performance)
 - MAPE < 2.8% (forecast error rate)
 - R² > 0.85 (explanation of variance)

- Business Impact:
- 15% reduction in average passenger wait time
 - 10% increase in driver utilization rate
 - 12% improvement in revenue per trip during peaks
 - 95% API uptime and <200ms response time

- Operational Excellence:
- Weekly model retraining success rate > 98%
 - Data pipeline reliability > 99.5%
 - False alarm rate for monitoring < 2%
 - Mean time to recovery for issues < 30 minutes

EXPECTED OUTCOMES

- Short-term (3 months):
- Deployed production forecasting system
 - 10-15% improvement in operational efficiency
 - Reduced customer complaints about wait times
 - Data-driven decision making for dispatch

- Medium-term (6-12 months):
- Advanced features and external data integration
 - Expansion to other cities/regions
 - Integration with third-party services
 - Significant competitive advantage in market

- Long-term (1+ years):
- Industry-leading prediction accuracy
 - Fully autonomous demand-supply optimization
 - Platform for additional ML/AI services
 - Foundation for autonomous vehicle integration

Report completed: September 12, 2025 at 10:42 PM